

1. Purpose

Define a distinctive, interactive, evidence-first interface for the candidate evaluation system. The UI must feel like a professional hiring investigation and decision workspace—not a generic AI dashboard, chatbot, or collection of four score cards.

2. Product Personality

Visual character: editorial, analytical, calm, human, evidence-driven, and slightly investigative. Avoid the stereotypical “AI SaaS” look: excessive gradients, glowing neural-network graphics, floating chatbot bubbles, generic robot icons, and endless cards with percentages.

3. Core UX Principle

The product should tell a story: SOURCE → PROFILE → INDEPENDENT OPINIONS → DEBATE → DECISION. Users should always understand what the system knows, what each agent thinks, where the evidence came from, and why the final recommendation was reached.

4. Primary User Journey

1. Create evaluation. 2. Upload resume and transcript. 3. Review extracted candidate profile. 4. Review evidence and requirement coverage. 5. Start evaluation. 6. Watch four independent agents work in parallel. 7. Explore each independent opinion. 8. Enter the debate timeline. 9. Watch disagreements and revisions emerge. 10. Review final decision. 11. Inspect unresolved disagreements and source evidence.

5. Information Architecture

Primary navigation: Evaluations, New Evaluation, Jobs, Candidates, Settings. Evaluation workspace navigation: Overview, Candidate Profile, Evidence, Agent Room, Debate, Decision, Report. Keep navigation shallow. The evaluation itself is the main workspace.

6. Evaluation Workspace

Use a persistent evaluation header containing candidate name, target role, pipeline state, and a compact progress indicator. The main content area changes by stage. Avoid making every section a dashboard grid.

7. New Evaluation Screen

Design this as a focused intake workspace rather than a form-heavy admin page. Show role selection, candidate identification, resume upload, transcript upload, and optional job description. Provide clear source-status indicators and an explicit “Build Candidate Profile” action.

8. Upload Interaction

Use drag-and-drop document zones with strong source identity. After upload, show filename, document type, file size, processing state, and extracted-page count when available. Provide a visible source preview or quick-open action.

9. Candidate Profile Screen

Present the profile as a structured candidate dossier. Sections: summary facts, skills, experience, projects, education, claims, technologies. Every important item should expose its evidence source. Distinguish visually between FACT, CANDIDATE CLAIM, INFERENCE, and UNKNOWN.

10. Evidence Interaction

Evidence should feel like a source citation, not a hidden tooltip. Clicking an Evidence ID opens a side panel containing source document, page/section, exact quote, validation status, and the finding(s) that use it. Preserve a clear path back to the current conclusion.

11. Requirement Coverage

Show job requirements alongside evidence coverage. Use states such as Supported, Partially Supported, Not Supported, and Unknown. Do not reduce the entire role fit to a single percentage.

12. Independent Agent Room

Show four distinct evaluator perspectives. Each persona gets a visual identity through typography, iconography, label, and interaction pattern—not cartoon avatars. The four opinions should be clearly separated to reinforce that they were generated independently.

13. Independence Visualization

The UI should explicitly communicate the architectural rule: “Independent evaluation — agents did not see each other's conclusions.” A simple visual sequence can show four parallel lanes converging only after the independent stage.

14. Agent Detail View

For each agent show: recommendation, confidence, assessment, strengths, concerns, uncertainties, requirement mappings, and evidence-backed findings. Make Evidence IDs clickable. Avoid leading with a giant score.

15. Finding Interaction

Each finding should be expandable. Collapsed state: finding title + severity + evidence count. Expanded state: statement, reasoning, evidence, source quote, and confidence. This keeps the page readable while preserving auditability.

16. Evaluation Progress

During processing, show a live stage tracker rather than a generic spinner. Stages: Documents → Profile → Independent Agents → Evidence Validation → Debate → Decision. For independent agents, show four parallel statuses.

17. Debate Experience

The debate should be the most interactive screen in the product. Present it as a chronological conversation between professional evaluators, not a chatbot. Each message should show agent identity, claim being challenged, response, evidence links, and whether the position changed.

18. Debate Message Design

A debate message should answer: WHO said it? WHAT point are they responding to? DO THEY agree/disagree? WHAT evidence supports the response? DID their position change? Include “Responds to D12” or equivalent traceability.

19. Disagreement Visualization

When agents disagree, surface the disagreement as a first-class object. Example: “Production-depth claim” → Technical: moderate confidence; Skeptic: high concern; Hiring Manager: low impact on hiring decision. Allow the user to open the underlying evidence.

20. Opinion Revision

When an agent changes its position, show a visible revision event. Display original recommendation/confidence, revised recommendation/confidence, reason for change, triggering debate message, and supporting evidence. Do not overwrite the original opinion visually.

21. Final Decision Screen

The decision screen should feel like a decision memo, not a scorecard. Lead with recommendation and confidence, then explain why. Follow with decisive factors, supporting evidence, concerns, resolved disagreements, unresolved disagreements, and an alternative recommendation that was considered.

22. Non-Averaging Explanation

Include a small “How this decision was reached” section. It should state that the final decision weighed role requirements, evidence strength, concerns, contradictions, debate outcomes, confidence, and remaining uncertainty. Never display a fake formula implying simple averaging.

23. Unresolved Disagreements

Give unresolved disagreements their own section. Each should show topic, agents involved, competing interpretations, evidence, and impact on the final decision. The UI must not hide disagreement just to make the result look confident.

24. Final Report

Provide a report view optimized for review and export. Recommended order: recommendation → confidence → executive rationale → role-fit factors → strengths → concerns → agent perspectives → debate summary → unresolved disagreements → evidence appendix.

25. Interactivity Requirements

Important interactions: expandable findings; clickable evidence; source side panels; requirement-to-evidence navigation; agent filtering; debate round filtering; jump-to-message; revision history; stage progress; hover/focus explanations; report section navigation.

26. Visual Language

Prefer strong typography, generous whitespace, thin rules, restrained surfaces, document-like sections, subtle status indicators, and clear hierarchy. Use color primarily for semantic states such as concern, support, uncertainty, and validated evidence—not decoration.

27. Avoid Generic AI Patterns

Do not use: giant hero text saying “AI-powered hiring”; generic chatbot UI; robot heads; glowing gradients; excessive glassmorphism; meaningless circular scores; twelve-card dashboards; decorative circuit-board backgrounds; animated AI particles; “magic” buttons.

28. Distinctive Concept

Recommended design metaphor: “Candidate Intelligence Dossier.” The application resembles a modern analyst’s case file where source documents, evidence, independent expert assessments, debate, and decision reasoning are connected. This metaphor should remain subtle and professional, not become a literal detective theme.

29. Responsive Design

Desktop is the primary target because the evaluation workspace benefits from wide layouts. Tablet should remain usable. Mobile should prioritize report reading, evidence inspection, and status tracking rather than attempting to reproduce the full desktop workspace.

30. Accessibility

Use semantic HTML, keyboard navigation, visible focus states, sufficient contrast, readable text sizes, accessible labels, reduced-motion support, and non-color indicators for statuses. Do not rely solely on color to distinguish agent opinions or concern levels.

31. Loading and Empty States

Loading states should explain the current pipeline stage. Empty states should explain what is missing and what action is available. Avoid generic “Loading...” screens. Example: “Transcript uploaded. Extracting source passages for evidence validation...”

32. Error UX

Errors should identify what failed, whether the evaluation can continue, and what the user can do. For an agent failure, show “Skeptic evaluation unavailable” rather than silently filling the space with a fabricated result.

33. Interaction Safety

Never present an unsupported LLM claim as verified evidence. Invalid evidence should have an explicit status. Candidate-provided prompt-injection text must never appear as an executable instruction. UI labels should distinguish source content from system reasoning.

34. Performance UX

Do not block the entire interface while one agent runs. Display completed agent results as they become available when policy permits, while maintaining the distinction between “partial progress” and “complete independent evaluation.” Debate should not be available until its prerequisites are satisfied.

35. Design System Components

Core reusable components: EvaluationHeader, StageTracker, DocumentSourceCard, ProfileItem, EvidenceChip, EvidenceDrawer, RequirementRow, AgentPanel, FindingAccordion, ConfidenceBadge, DebateTimeline, DebateMessage, DisagreementCard, RevisionEvent, DecisionMemo, DecisionFactor, UnresolvedDisagreement, ReportNavigator, StatusIndicator.

36. Suggested Screen Map

Screen 1: Evaluation list. Screen 2: New evaluation. Screen 3: Processing workspace. Screen 4: Candidate Profile. Screen 5: Evidence explorer. Screen 6: Agent Room. Screen 7: Agent detail. Screen 8: Debate. Screen 9: Final Decision. Screen 10: Full Report.

37. Frontend State Model

Track evaluation status, document status, profile status, agent statuses, debate status, decision status, selected evidence, selected agent, selected debate message, and report section. Prefer server state from the API as the source of truth. Do

not duplicate LLM reasoning in frontend state.

38. Animation

Use restrained transitions to communicate stage changes, panel expansion, message arrival, and revision events. Avoid continuous decorative animation. Debate messages may enter sequentially to make the process understandable, but the UI must remain usable with reduced motion.

39. Demo Experience

For a college demonstration, the first screen should quickly communicate the product's unique concept. A reviewer should be able to understand within seconds: this system compares four independent AI evaluators, lets them debate, and produces an evidence-traceable final decision.

40. UX Acceptance Criteria

The UI is successful when a reviewer can: upload both documents; identify what facts came from the sources; inspect evidence behind a claim; see four independent opinions; understand that independence happened before debate; follow a direct agent-to-agent disagreement; see a position change; understand why the final decision was reached without averaging; and identify unresolved disagreement.

41. Design Quality Gate

Before implementation is considered complete, review every screen for generic AI styling, information overload, unsupported visual claims, unclear source traceability, hidden disagreement, and confusing stage transitions. The interface should prioritize evidence and reasoning over decorative AI branding.